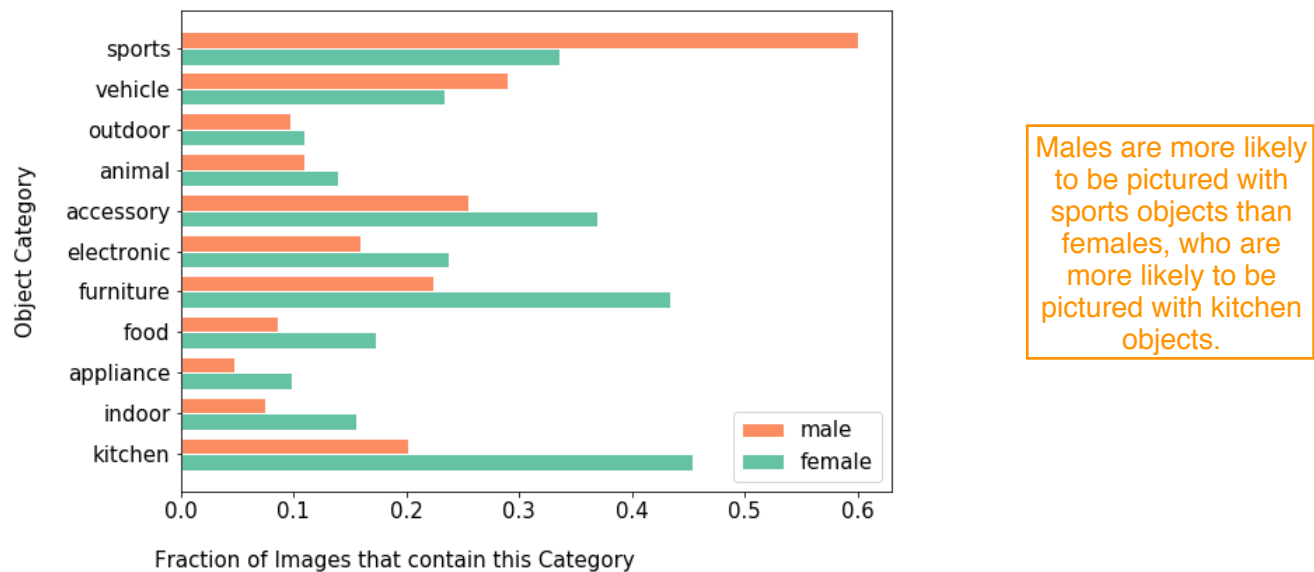



Gender-Based Summary

Overview Statistics

(M2) Distribution of object categories each gender appears with, sorted by ratio between the two.



(M11) Distribution of scenes that each gender appears in, sorted by ratio between the two.




Sample Interesting Findings


(M1) Probability image is labeled male when it should not be, i.e. given there's no face detected or person is too small: 0.7699

industrial and construction is the scene where the label of male is most likely to be picked over that of female

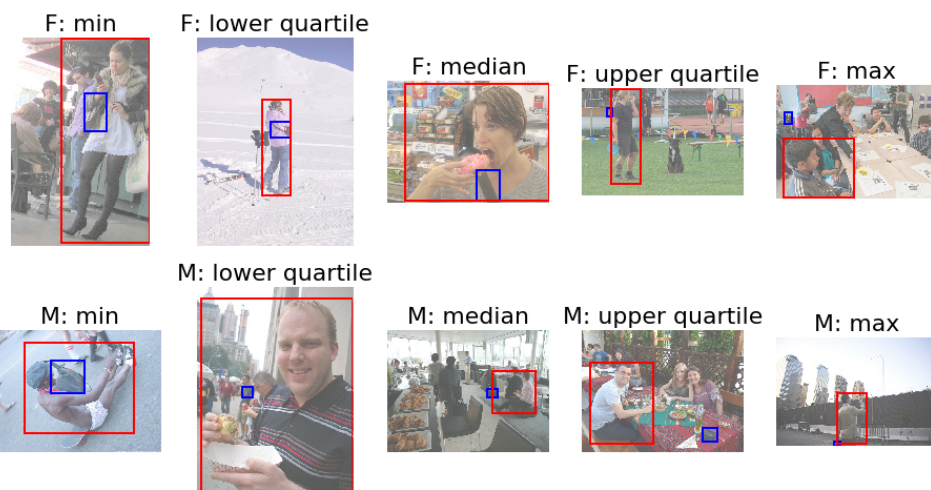
(M1) Fraction of image taken up by a person is different between the genders with a p-value of 2.597e-20, distribution shown below



(M2) Objects that are most statistically significantly represented with one gender over the other.



(M3) Qualitative example of handbag, which has the biggest ratio in distance between object and person (which can be interpreted as a proxy for interaction) between the genders. Females are closer than males. There is a red box around the person, and blue box around the object.



We can see that it does appear to be the case that when a male is pictured with a handbag, he is less likely to actually be interacting with it, whereas when a female is pictured with a handbag, she is actually interacting with it.

Some of the other metrics in the notebook

- (M2) Cooccurrence differences of objects between genders
- (M4) Scene differences per object between genders